

Bandwidth and Latency Measurement on Inter-Process Communications

Kedar Bastakoti, Ruohui Wang

Abstract

There are multiple ways to do inter-process communications(IPC). For example, pipe is often used to communicate between processes, whereas TCP and UDP can also be used to communicate both locally and remotely. Due to their differences in design, they tend to have different latencies and throughputs. Our experiments showed that the packet size is an important factor for both the latency and the throughput of IPC.

1 Clock Resolution

To begin with, how can we measure the time? We need clocks to do so. The Linux operating system, along with the underlying hardware, provides multiple methods for such purpose. But how can we choose one? One of the criteria is how precise they can be. And resolution is an important standard for the precision.

The resolution is the smallest possible increase of the clock. In order to measure the resolution, we tried to create minimal possible differences between two time measurements. We create such difference by inserting an line of assembly nop into the code.

As per the POSIX manual[1], the function `gettimeofday` is obsolete, and the switch to `clock_gettime` is recommended, so we used `clock_gettime` instead. along with `clock_gettime`, function `clock_getres` is provided for user to query the resolution of time. As we will demonstrated, it produces the same result as ours.

Also, the x86 CPU provides `rdtscp` instruction to give the CPU timestamp, in terms of TSC(Time Stamp Counter) cycles. As Sergiu and Terry pointed out, in order to use TSC as a reliable clocksource, it must be stable and have a constant rate, which can be examined using `\proc\cpuinfo[2]`.

2 System Calls

Since we are measuring I/O performance, we cannot avoid doing a lot of system calls. System calls involving context switching, and trapping into the kernel, both of which might . In order to get an better idea of how these system calls take up time. We measured some trivial system calls to see if the influence is non-trivial.

In order to measure the impact of trapping, we repeatedly make system calls that do least extra work. Here, we chose `getpid` and `getuid` as our syscalls. In addition, to avoid function returns that may take up time, we directly used the inline assembly to call the `syscall` instruction. And we return the minimal result in order to eliminate the effects random events including page faults and preemptions.

3 Pipe

We measured the performance of pipes in terms of latency and throughput. In order for a pipe to be used, we first create a pair of pipes, and then fork the process to let them talk to each other.

We tried to measure the impact of how the bytes written each time impacts the latency and the throughput. Since the pipe does not transmit messages in the form of a packet, we simulate it by initiating `read` and `write` each time with the desired size.

We measure the latency in terms of RTT. And repeat the experiment 100000 times to get the minimum latency as the final result. For latency, we make one process send 1GiB data, and the other process, on receiving all data, returns a message to the sender, and the time of the whole process is used to calculate the throughput. We apply the same method to the measurements on TCP/IP and UDP.

4 TCP

TCP uses sockets to communicate between the hosts which is used across local or remote communication. TCP uses byte-ordered reliable communication mechanism which is widely supported across today's operating systems. In order to measure the latency and throughput of the TCP stream sockets, packet sizes of 4, 16, 64, 256, 1K, 4K, 16K, 64K, 256K, and 512K were considered.

We hypothesized that the latency will proportionally increase with the packet size, and the throughput will also increase as the packet sizes are increased. To get the most accurate representation of the results the test for each individual transfer unit was repeated 1 million times and the minimum number amongst those 1 million trials was considered to be the value closest to the real value.

One of the issues encountered while measuring TCP latency was that when using the read and write function calls of the Rust language on TCP stream, the functions were reading and writing random bytes of data which were different than the expected values. This caused discrepancies in the read and write sequence of the server and the client model. To remedy this problem, we switched to different function calls, `read_to_end` and `write_all`, which helped synchronize the flow of data between the client and the server.

5 UDP

Similar to TCP, UDP also uses sockets for communication. However, unlike TCP, UDP doesn't provide reliability mechanism or any guarantees on the transmission of the data. We hypothesized that similar to UDP, the TCP latency should increase as packet sizes increase, and the throughput will also increase as the packet sizes are increased.


As UDP didn't have the builtin mechanism to guarantee whether the data were received at the other end or not, it caused issues while performing the latency and throughput tests in various stages. Unlike TCP, there is no concept of server and client. Instead, we have a port number associated with each process and both of the processes send and receive using the port number of the other processes. When the

sequence of read and write function calls were inconsistent the program would hang crash after a while.

Another error we encountered while write the benchmark was associated with how Rust's arrays work. A stack overflow error was received while running the throughput test. So, it was natural to check the function calls and the call stack to debug the issue. However, it was detected that the stack overflow occurred due to the allocation of the arrays on the stack. The array which was used to test the throughput was larger than 1 GiB. Even though, it was statically declared, it passed the compilation and only returned as a runtime error. To remedy this problem we allocated the array on heap using the Rust language's Vector type.

6 Evaluation

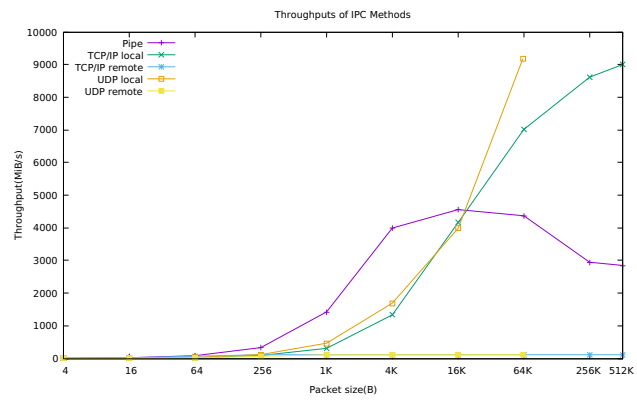
We ran the experiments on the instructional machines, with Intel(R) Core(TM) i5-4570 CPU @ 3.20GHz, with 8GB RAM. The results are followed.



7 Conclusion

References

- [1] IEEE, O. The open group base specifications issue 7, 2018 edition, 2017.



[2] SERGIU IORDACHE, T. L. TSC resynchronization, 2019.